

Problem 1: SPAM, SPAM, HAM

```
In [27]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import cvxopt
```

```
In [28]: spam_train = np.loadtxt('spam_train.data', delimiter=',')
spam_validation = np.loadtxt('spam_validation.data', delimiter=',')
spam_test = np.loadtxt('spam_test.data', delimiter=',')
```

```
In [29]: def preprocess(data):
    m, n = data.shape

    # X,y split
    X = data[:, :n-1]
    y = data[:, n-1:]

    # Set y = -1
    y = np.apply_along_axis(lambda x: -1 if x == 0 else 1, 1, y).reshape(-1, 1)

    return X,y
```

```
In [30]: X_train, y_train = preprocess(spam_train)
X_validate, y_validate = preprocess(spam_validation)
X_test, y_test = preprocess(spam_test)
```

```
In [31]: def get_accuracy(X, y, w, b,  $\lambda$ _ =None, y_ = None, X_ = None,  $\sigma$ 2=None):
    def gaussian_kernel(x, y,  $\sigma$ 2):
        return np.exp(-np.linalg.norm(x-y)**2 / (2 *  $\sigma$ 2))
    m, n = X.shape
    if w is not None:
        z = np.dot(X,w) + b
        f = (y * z) > 0
    else:
        print(f"Computing accuracy for  $\sigma$ 2 = { $\sigma$ 2}")
        y_predict = np.zeros(m)
        for i in range(m):
            wx = 0
            for  $\lambda$ , s1, sv in zip( $\lambda$ _, y_, X_):
                wx +=  $\lambda$  * s1 * gaussian_kernel(X[i], sv,  $\sigma$ 2)
            y_predict[i] = wx
        y_predict + b
        f = (y_predict * y.ravel()) > 0
    return np.sum(f.astype('float32')) * 100/m
```

```

In [32]: # Primal problem
def SVM_primal(X, Y, c):
    print(f"Computing c = {c}")
    m, n = X.shape

    P = np.zeros((m+n+1, m+n+1))
    P[:n,:n] = np.eye(n,n)
    P = cvxopt.matrix(P)

    q = np.zeros((m+n+1, 1))
    q[n:m+n,0] = c
    q = cvxopt.matrix(q)

    G = np.zeros((2*m, m+n+1))
    for i in range(m):
        for j in range(n):
            G[i][j] = -1 * Y[i] * X[i][j]
        G[i][n+i] = -1
        G[i][m+n] = -1 * Y[i]
        G[m+i][n+i] = -1
    G = cvxopt.matrix(G)

    h = np.zeros((2*m, 1))
    h[:m,0] = -1
    h = cvxopt.matrix(h)

    # CVXOPT Solver
    solution = cvxopt.solvers.qp(P, q, G, h)
    sol = np.array(solution['x'])
    w = sol[:n]
    b = sol[m+n]
    return w, b

# Dual problem
def SVM_dual(X, Y, c, σ2):
    print(f"Computing c = {c} and variance = {σ2}")
    m, n = X.shape

    # Create Gram matrix
    K = np.zeros((m, m))
    # Use gaussian kernel
    X_sq = -2 * np.dot(X, X.T)
    X_sq += (X ** 2).sum(axis=1).reshape(-1, 1)
    X_sq += (X ** 2).sum(axis=1)

    K = X_sq / (-2 * σ2)
    np.exp(K, K)
    # P = Combination of Yi Yj Xi Xj
    P = cvxopt.matrix(np.outer(Y,Y) * K)
    q = cvxopt.matrix(np.ones(m) * -1)

    # Constraints λY = 0
    A = cvxopt.matrix(Y, (1,m), 'd')
    b = cvxopt.matrix(0.0)

```

```

#  $\lambda \geq 0$ 
lhs = np.diag(np.ones(m) * -1)
lhs2 = np.identity(m)
G = cvxopt.matrix(np.vstack((lhs, lhs2)))
rhs = np.zeros(m)
rhs2 = np.ones(m) * c
h = cvxopt.matrix(np.hstack((rhs, rhs2)))

# CVXOPT Solver
solution = cvxopt.solvers.qp(P, q, G, h, A, b)

# Solver produces  $\lambda$ 
 $\lambda$  = np.ravel(solution['x'])
n_ $\lambda$  = len( $\lambda$ )
sv =  $\lambda$  > 1e-5
idx = np.arange(n_ $\lambda$ )[sv]
l =  $\lambda$ [sv]
support_labels = Y[sv]
support_vectors = X[sv]

b = 0.0
for n in range(len(l)):
    b += support_labels[n]
    b -= np.sum(l * support_labels * K[idx[n],sv])
b /= len(l)

return b, l, support_labels, support_vectors

```

```

In [33]: def run_svm(c_list, o2_list = [None], primal = True):
    data = {
        'c': [],
        'variance': [],
        'Training Data Accuracy': [],
        'Validation Data Accuracy': []
    }
    best_c, best_o2, best_validation_acc, best_train_acc = c_list[0], o2_list[0], 0, 0
    best_w, best_b = [], 0
    for c in c_list:
        for o2 in o2_list:
            data['c'].append(c)
            data['variance'].append(o2)
            w, b, λ, support_labels, sv = None, None, None, None, None
            if primal:
                w, b = SVM_primal(X_train, y_train, c)
                train_acc = get_accuracy(X_train, y_train, w, b)
                validation_acc = get_accuracy(X_validate, y_validate, w, b)
            else:
                w = None
                b, λ, support_labels, sv = SVM_dual(X_train, y_train, c, o2)
                train_acc = get_accuracy(X_train, y_train, None, b, λ_ = λ, y_ = y_train)
                validation_acc = get_accuracy(X_validate, y_validate, None, b, λ_ = λ, y_ = y_validate)
            data['Training Data Accuracy'].append(train_acc)
            data['Validation Data Accuracy'].append(validation_acc)
            if validation_acc > best_validation_acc or (validation_acc == best_validation_acc and train_acc > best_train_acc):
                best_validation_acc = validation_acc
                best_train_acc = train_acc
                best_c = c
                best_o2 = o2
                best_w = w
                best_b = b
                best_λ = λ
                best_sv = sv
                best_sv_y = support_labels
    df = pd.DataFrame.from_dict(data)
    if primal:
        test_acc = get_accuracy(X_test, y_test, best_w, best_b)
    else:
        test_acc = get_accuracy(X_train, y_train, None, best_b, λ_ = best_λ, y_ = y_train)
    df['Testing Data Accuracy'] = df.apply(lambda row: test_acc if row['c'] == best_c else None, axis=1)
    if best_o2 is None:
        df.drop(columns=['variance'], inplace=True)
    return df

```

1. Primal SVMs

- Using gradient descent or quadratic programming, apply the SVM with slack formulation to train a classifier for each choice of $c \in 1, 10, 10^2, 10^3, 10^4, 10^5, 10^6, 10^7, 10^8$ without using any feature maps.
- What is the accuracy of the learned classifier on the training set for each value of c ?
- Use the validation set to select the best value of c . What is the accuracy on the validation set for each value of c ?

- Report the accuracy on the test set for the selected classifier.

```
In [34]: c_list = [1,10,10**2,10**3,10**4,10**5,10**6,10**7,10**8]
df_primal = run_svm(c_list, primal=True)
```

Computing c = 1

	pcost	dcost	gap	pres	dres
0:	-1.8124e+03	1.0398e+04	8e+04	6e+00	4e+04
1:	5.6137e+03	-7.3332e+03	2e+04	1e+00	7e+03
2:	3.7452e+03	-2.4540e+03	7e+03	4e-01	2e+03
3:	2.3041e+03	-9.9333e+02	4e+03	2e-01	1e+03
4:	1.6237e+03	-3.8918e+02	2e+03	1e-01	6e+02
5:	1.3035e+03	-9.5546e+01	2e+03	6e-02	4e+02
6:	1.1640e+03	1.5815e+01	1e+03	4e-02	3e+02
7:	1.1113e+03	8.9810e+01	1e+03	3e-02	2e+02
8:	1.0293e+03	1.6990e+02	9e+02	2e-02	1e+02
9:	8.2001e+02	2.9315e+02	5e+02	1e-02	6e+01
10:	7.1530e+02	3.5051e+02	4e+02	6e-03	4e+01
11:	6.2597e+02	3.9397e+02	2e+02	3e-03	2e+01
12:	5.6753e+02	4.2284e+02	1e+02	2e-03	1e+01
13:	5.3370e+02	4.3939e+02	1e+02	1e-03	6e+00
14:	5.1168e+02	4.5038e+02	6e+01	5e-04	3e+00
15:	4.8730e+02	4.6134e+02	3e+01	3e-05	2e-01
16:	4.7777e+02	4.6801e+02	1e+01	7e-06	4e-02

```
In [35]: df_primal.style.apply(lambda x: ['background: lightgreen' if not np.isnan(x['Test
```

Out[35]:

	c	Training Data Accuracy	Validation Data Accuracy	Testing Data Accuracy
0	1	94.466667	93.500000	nan
1	10	94.733333	93.875000	nan
2	100	94.866667	93.875000	62.172285
3	1000	94.833333	93.750000	nan
4	10000	94.833333	93.750000	nan
5	100000	94.833333	93.750000	nan
6	1000000	94.833333	93.750000	nan
7	10000000	94.833333	93.750000	nan
8	100000000	94.833333	93.750000	nan

2. Dual SVMs with Gaussian Kernels

- Using quadratic programming, apply the dual of the SVM with slack formulation to train a classifier for each choice of $c \in 1, 10, 10^2, 10^3, 10^4, 10^5, 10^6, 10^7, 10^8$ using a Gaussian kernel with $\sigma^2 \in .1, 1, 10, 100, 1000$.
- What is the accuracy of the learned classifier on the training set for each pair of c and σ ?
- Use the validation set to select the best value of c and σ . What is the accuracy on the validation set for each pair of c and σ ?
- Report the accuracy on the test set for the selected classifier.

```
In [10]: c_list = [1,10,10**2,10**3,10**4,10**5,10**6,10**7,10**8]
σ2_list = [.1,1,10,100,1000]
df_dual = run_svm(c_list, σ2_list, primal=False)
```

Computing c = 1 and variance = 0.1

	pcost	dcost	gap	pres	dres
0:	-1.1526e+03	-5.7019e+03	1e+04	2e+00	4e-16
1:	-1.1153e+03	-3.2841e+03	2e+03	9e-13	2e-16
2:	-1.2019e+03	-1.4313e+03	2e+02	2e-12	1e-16
3:	-1.2565e+03	-1.2861e+03	3e+01	2e-12	1e-16
4:	-1.2645e+03	-1.2702e+03	6e+00	6e-12	6e-17
5:	-1.2662e+03	-1.2669e+03	8e-01	2e-12	6e-17
6:	-1.2663e+03	-1.2665e+03	2e-01	5e-12	5e-17
7:	-1.2663e+03	-1.2664e+03	4e-02	5e-13	5e-17
8:	-1.2663e+03	-1.2663e+03	1e-02	8e-12	6e-17
9:	-1.2663e+03	-1.2663e+03	5e-03	5e-12	7e-17
10:	-1.2663e+03	-1.2663e+03	8e-04	1e-12	5e-17

Optimal solution found.

Computing c = 1 and variance = 1

	pcost	dcost	gap	pres	dres
0:	-1.1106e+03	-5.6547e+03	1e+04	2e+00	6e-16
1:	-1.0753e+03	-3.2369e+03	2e+03	2e-13	3e-16
2:	-1.1581e+03	-1.3869e+03	2e+02	8e-13	2e-16

```
In [11]: df_dual.style.apply(lambda x: ['background: lightgreen' if not np.isnan(x['Testir
```

```
Out[11]:
```

	c	variance	Training Data Accuracy	Validation Data Accuracy	Testing Data Accuracy
0	1	0.100000	99.974000	19.623000	nan
1	1	1.000000	99.932000	26.389000	nan
2	1	10.000000	98.133333	68.755000	nan
3	1	100.000000	90.100000	79.625000	nan
4	1	1000.000000	83.233333	79.375000	nan
5	10	0.100000	99.973333	19.755000	nan
6	10	1.000000	99.973333	26.125000	nan
7	10	10.000000	99.773333	69.125000	nan
8	10	100.000000	97.736667	83.125000	nan
9	10	1000.000000	93.800000	85.875000	nan
10	100	0.100000	99.973333	19.125000	nan
11	100	1.000000	99.433333	26.125000	nan
12	100	10.000000	99.733333	69.125000	nan
13	100	100.000000	99.066667	84.125000	nan
14	100	1000.000000	97.033333	90.875000	nan
15	1000	0.100000	99.433333	19.125000	nan
16	1000	1.000000	99.433333	26.125000	nan
17	1000	10.000000	99.900000	68.125000	nan
18	1000	100.000000	99.966667	84.125000	nan
19	1000	1000.000000	98.900000	91.375000	nan
20	10000	0.100000	99.433333	19.125000	nan
21	10000	1.000000	99.433333	26.125000	nan
22	10000	10.000000	99.900000	68.125000	nan
23	10000	100.000000	99.533333	84.125000	nan
24	10000	1000.000000	99.000000	92.125000	81.125000
25	100000	0.100000	99.433333	19.125000	nan
26	100000	1.000000	99.433333	26.125000	nan
27	100000	10.000000	99.900000	68.125000	nan
28	100000	100.000000	99.866667	84.250000	nan
29	100000	1000.000000	99.433333	91.125000	nan
30	1000000	0.100000	99.433333	19.125000	nan
31	1000000	1.000000	99.433333	26.125000	nan
32	1000000	10.000000	99.900000	67.125000	nan
33	1000000	100.000000	99.533333	73.000000	nan

	c	variance	Training Data Accuracy	Validation Data Accuracy	Testing Data Accuracy
34	1000000	1000.000000	89.366667	72.125000	nan
35	10000000	0.100000	99.433333	19.125000	nan
36	10000000	1.000000	99.433333	26.125000	nan
37	10000000	10.000000	99.900000	68.125000	nan
38	10000000	100.000000	84.533333	69.000000	nan
39	10000000	1000.000000	87.200000	70.125000	nan
40	100000000	0.100000	99.433333	19.125000	nan
41	100000000	1.000000	99.433333	26.125000	nan
42	100000000	10.000000	99.900000	68.125000	nan
43	100000000	100.000000	99.533333	71.000000	nan
44	100000000	1000.000000	81.266667	53.125000	nan

3. k-Nearest Neighbors

- What is the accuracy of the k-nearest neighbor classifier for k = 1,5,11,15,21?


```
In [33]: from collections import defaultdict
```

```
class kNN:
```

```
    def __init__(self, neighbors=1):
```

```
        self.train_data = []
```

```
        self.labels = []
```

```
        self.m = 0
```

```
        self.n = 0
```

```
        self.k = neighbors
```

```
    def fit(self, X_train, y_train):
```

```
        self.train_mean = X_train.mean(axis=0)
```

```
        self.train_std = X_train.std(axis=0)
```

```
        self.train_data = X_train
```

```
        self.train_data_normed = (self.train_data - self.train_mean) / self.train_std
```

```
        self.m, self.n = self.train_data.shape
```

```
        self.labels = y_train
```

```
    def get_distances(self, X_test):
```

```
        distances = -2 * self.train_data_normed.dot(X_test.T) + np.sum(X_test**2, axis=1)
```

```
        distances[distances < 0] = 0
```

```
        return distances
```

```
    def predict(self, X_test):
```

```
        X_test_normed = (X_test - self.train_mean) / self.train_std
```

```
        distances = self.get_distances(X_test_normed)
```

```
        idx = np.argsort(distances, axis=0)
```

```
        idx = idx[0:self.k, :]
```

```
        m, n = idx.shape
```

```
        labels = self.labels.ravel()
```

```
        y_pred = np.zeros((X_test.shape[0], 1))
```

```
        for col in range(n):
```

```
            classes = defaultdict(int)
```

```
            for row in range(m):
```

```
                label = labels[idx[row, col]]
```

```
                classes[label] += 1
```

```
            # Get the majority class
```

```
            y_pred[col] = max(classes, key=classes.get)
```

```
        return y_pred
```

```
    @staticmethod
```

```
    def get_accuracy(y_pred, y_test):
```

```
        return np.mean(y_pred.flatten() == y_test.flatten()) * 100
```

```
In [34]: def run_kNN(k_list):
    data = {
        'k': [],
        'Validation Accuracy': [],
        'Test data Accuracy': []
    }
    for k in k_list:
        data['k'].append(k)
        classifier = kNN(k)
        classifier.fit(X_train, y_train)
        y_validate_pred = classifier.predict(X_validate)
        data['Validation Accuracy'].append(kNN.get_accuracy(y_validate_pred, y_validate))
        y_test_pred = classifier.predict(X_test)
        data['Test data Accuracy'].append(kNN.get_accuracy(y_test_pred, y_test))

    return pd.DataFrame.from_dict(data)
```

```
In [35]: k_list = [1,5,11,15,21]
df = run_kNN(k_list)
display(df)
```

	k	Validation Accuracy	Test data Accuracy
0	1	88.500	72.659176
1	5	89.625	70.786517
2	11	88.500	72.159800
3	15	86.375	71.535581
4	21	86.500	70.911361

```
In [10]: m, n = X_train.shape
X = X_train
c = 10
Y = y_train
P = np.zeros((m+n+1, m+n+1))
P[:n,:n] = np.eye(n,n)
q = np.zeros((m+n+1, 1))
q[n:m+n,0] = c
G = np.zeros((2*m, m+n+1))
for i in range(m):
    for j in range(n):
        G[i][j] = -1 * Y[i] * X[i][j]
    G[i][n+i] = -1
    G[i][m+n] = -1 * Y[i]
    G[m+i][n+i] = -1

h = np.zeros((2*m, 1))
h[:m,0] = -1
```

```
In [26]: h2 = np.zeros((m*2, 1))  
h2[:m,0] = -1  
np.mean(h2.ravel() == h.ravel())
```

```
Out[26]: 1.0
```

4. Which of these approaches (if any) should be preferred for this classification task?

Explain

SVM with gaussian kernel should be preferred for this classification task as it has higher test data accuracy of 90.38% compared to kNN model which has average test data accuracy ~70%. SVM performs better in Higher dimensions compared to kNN.